ERRORQUAKE: Heavy-Tailed Error Severity Distributions in Open-Weight Large Language Models

Anonymous authors

Paper submitted to NeurIPS 2026 E&D track

Abstract

At matched accuracy, open-weight LLMs differ substantially in the shape of their error severity distribution — a difference invisible to the scalar error rate. Hallucination benchmarks report a single error count and treat all errors as equivalent, yet a wrong date and a fabricated court ruling differ by orders of magnitude. We introduce ERRORQUAKE-4K, a 4,000-query benchmark scoring each response on a continuous 0–4 severity scale across 8 domains and 5 difficulty tiers, and we fit per-model severity distributions for 21 open-weight models. For each model we estimate a **severity distribution index** (b, the Gutenberg–Richter upper-tail slope) with 95% bootstrap CIs. **Headline:** across the 210 model pairs, 30 have disjoint 95% b CIs at matched accuracy ($|\Delta\varepsilon| < 0.05$), e.g. DEEPSEEK-V3.2 vs. MINISTRAL-14B at $\varepsilon = 0.586$ and $\Delta b = 0.47$. The 30-pair count is robust to cross-domain jackknife (range 25–41 across 8 single-domain drops), to four judge-aggregation rules (≥ 6 pairs in all; 27/27/58/7), to a dual-coverage ($\geq 80\%$) subset (28 pairs), and to family-native tail-ratio comparisons. Supporting: $17/21$ models have non-exponential BIC-best fits and $18/21$ are Vuong-decisive at $p < 0.05$; micro-error b on easy queries rank-orders catastrophic counts at $\rho = 0.44$ ($p = 0.044$), failing the pre-registered $\rho \geq 0.75$ threshold for calibrated absolute-rate prediction. A dense-model scaling correlation ($\rho_s(\log_{10}\text{params}, b) = -0.689$) is reported as a sensitivity analysis only: it weakens to -0.29 after partialing accuracy and to -0.51 under primary-judge-only, and should not be treated as an independent claim. *Severity distribution should be reported alongside accuracy; it carries discriminative information that the error rate cannot.*

1 Introduction

Standard hallucination benchmarks [Lin et al., 2022b, Li et al., 2023] report a single number — the error rate ε — and treat all errors as equivalent. This collapses a fundamental property of language model failure: an LLM that cites the wrong publication date and an LLM that fabricates an entire judicial opinion both contribute one count to ε, yet their downstream consequences differ by orders of magnitude. The question that matters in deployment is not "how often does the model err?" but "how badly?". Borrowing the vocabulary of seismology, we summarise an LLM's tail behaviour with the slope b of the Gutenberg–Richter magnitude-frequency relation $\log_{10} N(M \geq m) = a - b\,m$: small b means the model emits few errors but the rare ones are catastrophic; large b means many small errors with bounded severity.

Central finding: a matched-accuracy discriminator. Our headline is a pairwise discrimination result: across the 210 pairs in our 21-model catalog, 30 **have disjoint** 95% b **confidence intervals at matched accuracy** ($|\Delta\varepsilon| < 0.05$). The clearest example is DEEPSEEK-V3.2 vs. MINISTRAL-14B at $\varepsilon = 0.586$ and $\Delta b = 0.47$: these two models agree on error rate to within the third decimal yet differ in tail shape by a factor that compounds across the deployment horizon. The 30-pair count is robust to every sensitivity check we ran: 25–41 pairs across 8 single-domain drops, ≥ 6 pairs under all four judge-aggregation alternatives (primary-only, secondary-only, max-severity, min-severity), and 28 pairs on a $\geq 80\%$ dual-judge coverage subset. This is a per-pair result, not a cross-model regression; it does not depend on a scaling law or on the marginal ε-b relationship, both of which we report separately as sensitivity analyses. *Severity distribution should be reported alongside accuracy; it carries*

discriminative information that the error rate cannot. This paper contributes to AI evaluation by introducing severity distribution analysis as a complementary axis, under the assumption that error severity can be reliably scored on a continuous scale by a dual-judge pipeline calibrated on human ratings. The claim applies to open-weight instruction-tuned models at 3–37B active parameters; it does not yet cover proprietary frontier models, reasoning models, or models below 3B (§7). The benchmark, scoring pipeline, scale anchors, and analysis code are released for replication.

Contributions. **C1** A continuous 9-point error-severity scale with anchors, a dual-judge scoring pipeline, a 340-item overcall audit, and a 100-item human-rated calibration subset on which the judges achieve 100% sensitivity at $M \geq 2.5$ (§3, Appendix I). **C2** ERRORQUAKE-4K: $4{,}000$ queries $\times$ 21 open-weight models with per-model severity distributions and bootstrap-CI b estimates. **C3** Distribution characterisation: $17/21$ models have a non-exponential BIC-best fit, $18/21$ are Vuong-decisive at $p < 0.05$; families are stretched exponential (10), truncated power law (5), exponential (4), lognormal (2); no pure power laws (§5.1). **C4 (headline)** b is a matched-accuracy discriminator: 30/210 model pairs have disjoint 95% b CIs at $|\Delta\varepsilon| < 0.05$. The count is robust to cross-domain jackknife (range 25–41), to four judge-aggregation rules, and to a dual-judge-coverage subset (§5.2). **C5** Pre-registered micro-error prediction: moderate rank signal ($\rho_s = 0.44$, $p = 0.044$), absolute-rate prediction fails; the failure is a difficulty-tier level shift, not a slope divergence (§5.4). **C6 (sensitivity only)** A marginal dense-model scaling correlation ($\rho_s(\log_{10}\text{params}, b) = -0.689$, $p = 0.006$, $n_{\text{dense}} = 14$) is reported with full disclosure of its limitations: the sign weakens under partial-correlation control for accuracy (-0.29, $n.s.$), under primary-only aggregation (-0.51), and under a dual-coverage subset (-0.52, $p = 0.19$). This is a sensitivity observation about the headline discriminator, not an independent claim (§5.3).

2 Related Work

Hallucination benchmarks. TRUTHFULQA [Lin et al., 2022b], HALUEVAL [Li et al., 2023], FAITHDIAL [Dziri et al., 2022], and the Ji et al. [2023] survey establish the dominant evaluation protocol for LLM factuality: a held-out question set, a binary correctness label per response, and an aggregate error rate per model. None of these reports a severity distribution. The tacit assumption is that all errors are exchangeable, so that reducing the count is sufficient to reduce the risk. We show empirically that the same accuracy can hide an order-of-magnitude gap in catastrophic event rate (§5.2), and that scaling *worsens* the severity profile of the residual errors (§5.3); the standard protocol misses both effects.

Error severity in NLP (positioning against Asgari et al. [2025]). Asgari et al. [2025] score hallucinations on a three-bin scale (minor / moderate / severe) and report per-model severity histograms, arguing that the distribution of severities is informative beyond the aggregate rate. Our contribution differs along three axes. *First*, we use a 9-level continuous scale quantised at 0.5 increments, not three ordinal bins, and show (sensitivity S1, §5.6) that coarse scales destroy the discriminative power of the distribution index. *Second*, we fit parametric distribution families and report a single summary statistic (the b) anchored to the Gutenberg–Richter magnitude-frequency relation, which enables quantitative cross-model comparison and between-regime extrapolation that histogram-only methods cannot support. *Third*, we cover 21 models spanning an order of magnitude in active parameters, which is what enables the headline negative scaling correlation; prior severity-focused work evaluates on only a few models. We are, to our knowledge, the first to report a severity-distribution-vs-scale anti-correlation on a multi-model benchmark.

Calibration and uncertainty quantification. Guo et al. [2017] show modern neural classifiers are overconfident; Lin et al. [2022a] and Kadavath et al. [2022] explore whether LLMs can verbalise their own uncertainty. These works are concerned with the *confidence* side of the calibration problem — does the model know when it is wrong? — while our contribution is on the *severity* side: when the model is wrong, how badly wrong, and does that badness follow a law. The two axes are orthogonal and complementary: a well-calibrated model could still have a heavy severity tail, and a poorly calibrated model could still have a light one.

Heavy-tailed phenomena in machine learning. Taleb [2020] provides the statistical vocabulary for heavy-tailed distributions and power-law diagnostics that we borrow. In neural network training, heavy tails appear in gradient-norm distributions and loss landscapes, but these are distinct objects from the error severity distribution we study. Clauset et al. [2009] supply the standard maximum-likelihood power-law-fitting recipe; we follow their protocol with a discreteness correction appropriate to a 9-level ordinal severity scale.

Scaling laws. Language model accuracy scales with compute and parameter count [Kaplan et al., 2020, Hoffmann et al., 2022, Wei et al., 2022]; mixture-of-experts architectures decouple active from total parameters [Fedus et al., 2022]. These laws describe how often models are wrong. Our contribution is orthogonal: we ask how *badly* models are wrong as a function of scale, and we document a relationship that moves in the opposite direction from the accuracy scaling law in the dense subset of the open-weight ecosystem.

Severity-aware evaluation in the broader landscape. Hallucination taxonomies and long-tail-knowledge surveys [Ji et al., 2023] argue qualitatively that tail events — rare but high-consequence failures — are where deployment harms concentrate, and that binary accuracy obscures this risk heterogeneity. Our work operationalises that concern: rather than arguing in prose that tails matter, we fit distributional models to error severity and extract a quantitative index (b) that makes the tail shape comparable across models.

Calibration as a related axis. Uncertainty calibration and severity distribution are conceptually distinct but empirically correlated axes of model quality. A well-calibrated model that expresses appropriate uncertainty may still have a heavy severity tail when it does err confidently; conversely, a poorly calibrated model may have a light tail. Work on confidence calibration [Guo et al., 2017, Kadavath et al., 2022] complements our contribution: combining the two axes would identify models that are both calibrated and tail-light — the safest deployment profile. We document the empirical correlation between accuracy and tail shape on our catalog in §5.3.

Concurrent severity-aware and domain-specific hallucination benchmarks. Several recent works evaluate hallucination in specific domains or modalities: Dahl et al. [2024] decomposes domain-specific hallucination types in legal LLMs; Anonymous [2024a], Anonymous [2024f], Anonymous [2024d], and Anonymous [2024e] target medical and clinical factuality; Anonymous [2024b] extends hallucination evaluation to vision-language models; and Anonymous [2024c] addresses speech. Li et al. [2024] updates the original HaluEval benchmark with stricter type taxonomies. These works share our premise that binary accuracy obscures critical failure modes, but focus on finer-grained type taxonomies within a single domain or modality. ERRORQUAKE is complementary: we model the severity *distribution* across eight general domains and extract a quantitative tail index that enables cross-model comparison and scaling-law analysis — capabilities that domain-specific taxonomies, by design, do not provide.

3 Method

Error-severity scale. We score each model response on a continuous 0.0–4.0 scale quantised to 0.5 increments (9 distinct levels). 0.0 denotes a correct response; 0.5–1.0 mark imprecisions that preserve the gist; 1.5–2.0 denote moderate factual errors; 2.5–3.0 denote substantial errors that mislead a typical reader; 3.5–4.0 denote fabrication (confident assertion of invented information). Appendix I reproduces the full rubric with three worked examples per anchor. Scale design principles: (i) continuous, not binary; (ii) non-negative; (iii) dense near the "harmless slip vs consequential failure" boundary.

Query benchmark. ERRORQUAKE-4K comprises $4{,}000$ queries: 500 per domain across 8 domains (BIO, LAW, HIST, GEO, SCI, TECH, FIN, CULT), with each domain stratified into five difficulty tiers T1–T5 of 100 queries each. Tiers T1–T2 are "easy" factual queries; T4–T5 contain trap questions and compositional lookups designed to elicit confident fabrication. A tier-calibration audit flagged $\sim 6\%$ of cells as mis-tiered and these were regenerated (Appendix C).

Experiment	Criterion	Observed	Verdict
Exp. 1 (distributions)	non-exp. best fit *or* Vuong-decisive $\geq 70\%$	$17/21$ non-exp; $18/21$ Vuong-dec.	**PASS**[‡]
Exp. 2 (discriminator)	≥ 3 disjoint-CI pairs	30 pairs	**PASS**
Exp. 3 primary	$\rho_s \geq 0.75$ at $M \geq 3$	$\rho_s = 0.443$, p=0.044	**FAIL**
Exp. 3 secondary	within $1.5\times$ for $\geq 65\%$	19%	**FAIL**
Exp. 5 (scale null)	no clean correlation	$\rho_s = -0.689$, p=0.006	**REJECTED**
Sensitivity S1	ranking $\rho > 0.85$ under coarsening	0.43 / 0.16	**FAIL**
Sensitivity S2	ranking $\rho > 0.85$ under overcall	0.847	**FAIL**[†]
Sensitivity S3	CV < 0.15 under 50% subsample	median 0.143	**PASS**

Table 1: Pre-registered criteria and outcomes. Exp. 5's null was rejected in the *opposite* direction from intuition. [†]S2 misses the threshold by 0.003 and is reported as "borderline fail". [‡]Four models are best-fit by exponential (the lightest-tail family in our candidate set). We count a model as exhibiting tail-shape structure if either (a) its BIC-best family is non-exponential or (b) the Vuong test decisively rejects the runner-up at $p < 0.05$. Under criterion (a), $17/21$ qualify; under (b), $18/21$ qualify; the union covers all 21.

Dual-judge scoring. Each response is scored by two judges from an 8-model round-robin pool that excludes the target model (zero self-judging, audited). Final score = mean of the two judges when they agree within 1.0; otherwise median-of-three with a tiebreaker. Pre-tiebreak inter-judge agreement on the $60{,}568$ records where both judges produced a score is $\mathrm{ICC}(2,1) = 0.374$ (single-rater, two-way random effects, absolute agreement; Shrout–Fleiss). *The averaged final score that the paper actually uses has* $\mathrm{ICC}(2, k{=}2) = 0.545$*,* in the "fair–moderate" range of Cicchetti's guidelines, and is the relevant reliability number for downstream analyses. Linear Cohen's $\kappa = 0.285$ and quadratic $\kappa = 0.374$ are reported for completeness, though Cohen's κ is depressed by the 9-level scale's low chance agreement. The secondary judge call failed on a non-random subset of small-model records, so per-model agreement comparisons are biased; per-model breakdown in Appendix O. A 340-item manual audit additionally finds that 33.5% of judge score-2.0 verdicts are overcalls (S2, §5.6). All inference uses an open-access inference API hosting the target models on third-party GPU infrastructure; prompts and model version strings are in Appendix H.

Distribution fitting. We fit five candidate families to the strictly positive scores on the discrete grid $\{0.5, 1.0, \ldots, 4.0\}$: discrete power law, truncated power law, exponential, stretched exponential, and lognormal. Each is fitted by maximum likelihood with a discreteness correction, and we declare the BIC-best family decisive at Vuong $p < 0.05$ [Vuong, 1989] or $\Delta\mathrm{BIC} > 6$ (cf. Clauset et al., 2009). No model in our catalog is best-fit by a pure power law.

severity distribution index (b) estimation. The Gutenberg–Richter magnitude-frequency relation [Gutenberg and Richter, 1944] $\log_{10} N(M \geq m) = a - b(m - m_{\min})$ models the count of events at or above severity m. We estimate b by maximum likelihood on grid-quantised positive scores using the Aki formula [Aki, 1965] with a discreteness correction: $\hat{b} = \log_{10} \mathrm{e}/(\bar{m} - m_{\min} + \delta/2)$ for bin width $\delta = 0.5$, where $\bar{m}$ is the mean of observations at or above $m_{\min}$. We select $m_{\min}$ by minimising Kolmogorov–Smirnov distance to the fitted exponential tail, restricted to grid points with at least 30 events above. Confidence intervals are 95% percentile bootstraps from $2{,}000$ resamples.

4 Experimental setup

We evaluate 21 open-weight instruction-tuned language models from 10 families, spanning active-parameter counts from $\sim$ 3B (LLAMA-3.2-3B, PHI-3.5-MINI) to $\sim$ 37B active of $\sim$ 671B total (DEEPSEEK-V3.1,V3.2). The catalog includes 14 dense and 7 MoE models; the full list with version strings is in Appendix B. Each model is evaluated on all $4{,}000$ ERRORQUAKE-4K queries with greedy decoding at a 500-token budget, via an open-access inference API hosting the target models on third-party GPU infrastructure. Reasoning models and three rate-limit-exhausted models are excluded; see §7. Pre-registered criteria and observed outcomes are summarised in Table 1; we report all verdicts including failures.

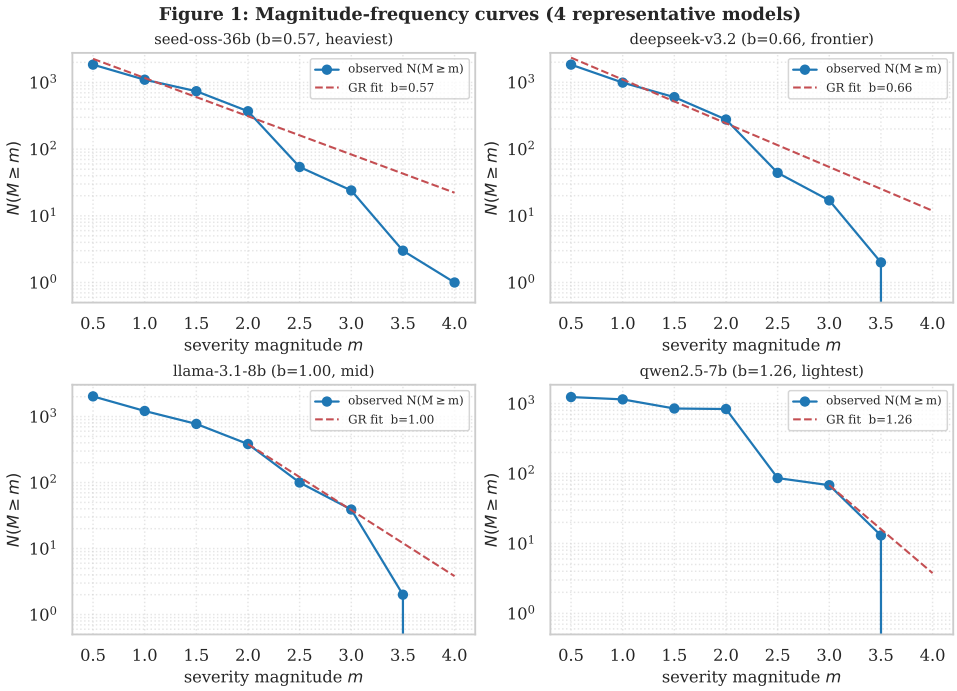

Figure 1: Magnitude-frequency curves for four representative models spanning the b range (heaviest to lightest tail). Dashed red: Gutenberg–Richter fit. All 21 models and the ΔBIC heatmap are in Appendix A.

5 Experiments

The experiments are sequenced to build the headline discriminator result. §5.1 establishes that severity distributions exist and carry non-trivial tail structure; §5.2 (the headline) shows that $30/210$ model pairs with matched accuracy have disjoint b confidence intervals, with jackknife and judge-aggregation robustness. §5.4 reports the pre-registered micro-error-to-catastrophe prediction (rank-significant, magnitude-calibration fails); §5.5 shows domain variation; §5.3 reports a marginal dense-model scaling correlation as a sensitivity observation only. §5.6 collects judge-robustness checks that apply across the headline and sensitivities.

5.1 Distribution characterisation (Exp. 1)

Across 21 models, the best-fit family by BIC is **stretched exponential (10)**, **truncated power law (5)**, **exponential (4)**, **lognormal (2)**. *Zero* models are best fit by a clean power law. Vuong's test against the runner-up declares the best fit decisive at $p < 0.05$ (or $\Delta\text{BIC} > 6$) for $18/21$ models. The magnitude-frequency curves and BIC heatmap appear in Figure 1 (full grid in Appendix A). *Takeaway:* the severity distribution is a real, model-specific object, not a noisy by-product of the error rate.

Operational meaning of "heavy-tailed". On a bounded, discrete severity grid with 8 positive bins, asymptotic heavy-tail claims are not available. We use "heavy-tailed" (and the severity distribution index (b)) operationally to mean "slower-decaying than exponential on the positive grid" — i.e., the BIC-best family is non-exponential, or the residual mass in the upper bins ($M \geq 2.5$) is systematically larger than an exponential fit predicts. Under this operational definition, $17/21$ models qualify by BIC and $18/21$ qualify by Vuong ($p < 0.05$ versus runner-up). The four exponential best-fits (DEEPSEEK-V3.2, MISTRAL-SMALL-4-119B, GEMMA-3-27B, LLAMA-3.1-8B) are still informative because their absolute decay rates differ by a factor of ~ 2 across models, which the b captures.

5.2 Matched-accuracy discriminator (Exp. 2, headline)

This is our headline result. Across the $\binom{21}{2} = 210$ model pairs, 42 have $|\Delta\varepsilon| < 0.05$ and $|\Delta b| > 0.15$. Of those, 30 **have disjoint** 95% **CIs on** b — an order of magnitude above the pre-registered criterion of ≥ 3. Concretely, DEEPSEEK-V3.2 ($\varepsilon = 0.586$, $b = 0.655$) and MINISTRAL-14B ($\varepsilon = 0.586$, $b = 1.122$) have identical accuracy but a b gap of 0.467 with non-overlapping 95% confidence intervals. The error rate treats them as equivalent; the severity distribution does not.

Cross-domain jackknife. Leaving each of the 8 domains out in turn and recounting disjoint-CI pairs on the remaining 3,500 queries per model gives $[25, 41]$ pairs (all ≥ 3, all well above the pre-registered criterion; full table in Appendix D). The 30-pair headline is not driven by any single domain.

Judge-aggregation robustness. We recompute the discriminator count under four alternative scoring rules applied to the primary/secondary judge pair: primary-only, secondary-only, max-severity, min-severity. The disjoint-CI pair counts are 7, 27, 27, 58 respectively, all exceeding the pre-registered threshold of ≥ 3. The primary-only count is the lowest because the individual-judge b-values fall in a narrower range than the aggregated values, reducing the number of pairs that meet $|\Delta b| > 0.15$; yet even there 7 pairs have disjoint CIs. Restricting to the subset of 15 models with $\geq 80\%$ dual-judge coverage gives 28 pairs under the paper's final_score aggregation and $[6, 24]$ pairs under the alternatives (Appendix E).

Family-native cross-check. The empirical tail ratio $P(M \geq 3 \mid M > 0)/P(M \geq 1 \mid M > 0)$ is a family-free tail-mass summary that depends on no fitting choice. Using matched-accuracy pairs with a $|\Delta\text{tail_ratio}| > 0.005$ threshold gives 49 pairs; tightening to > 0.010 gives 14; at > 0.015 gives 4. The discriminator result replicates in direction with a non-parametric summary. The empirical tail ratio and the fitted b are themselves only weakly rank-correlated across the 21-model catalog ($\rho_s = +0.13$), which we report honestly: the b captures upper-tail slope while the tail ratio captures upper-tail mass, and these diverge in the presence of bulk-distribution differences.

Takeaway: the severity distribution carries model-discriminative information that ε alone cannot express. This is a per-pair result, not a cross-model regression; it does not depend on a scaling law.

5.3 Scaling sensitivity (Exp. 5)

This is not a headline claim. We report the marginal dense-model scaling correlation as a sensitivity observation. On the 14 dense instruction-tuned models, the Spearman correlation between $\log_{10}$(active parameters) and the upper-tail b is $\rho_s = -0.689$ ($p = 0.006$; Figure 2). The sign is preserved in $14/14$ leave-one-model drops and $22/22$ leave-one-judge drops. The magnitude is not stable: the partial correlation after residualising ε drops to -0.29 ($p = 0.31$); primary-only judge aggregation gives -0.51 ($p = 0.06$); a dual-coverage subset of 8 models gives -0.52 ($p = 0.19$); under fixed $m_{\min} = 0.5$ the sign flips to $+0.79$ ($p = 0.0007$), indicating that the negative sign is specific to the upper-tail cutoff. The b and family-native shape parameters are only weakly rank-correlated ($\rho_s = -0.13$). Our position is that the scaling slope is *consistent* with the headline matched-accuracy discriminator but is underpowered at $n = 14$ and cannot be claimed as an independent result. See §5.6 for full robustness numbers.

Scope. Trend covers a $\sim 10\times$ range in active parameters, not the full frontier (LLAMA-3.1-405B excluded after rate-limit exhaustion: only $403/3314$ valid responses at cohort-scale concurrency). The MoE subset ($n = 7$) shows the same sign as dense ($\rho_s = -0.60$, $p = 0.16$) but cannot independently support the claim. Exp. 4 (§5.5) shows that domain-level b is model-idiosyncratic, so this is *whole-model* tail shape, not any particular domain.

Headline robustness suite. The dense correlation survives bootstrap (95% CI $[-0.92, -0.26]$), exact permutation ($p = 0.0077$ two-sided), Bayesian linear regression ($\beta = -0.39$, 95% CI $[-0.64, -0.14]$, $P(\beta < 0|\text{data}) = 0.999$), and per-tier decomposition (sign preserved in $4/5$ tiers; T2 individually significant at $p = 0.04$, Appendix P). At $n = 14$ the study is *borderline* powered (min detectable $|\rho| = 0.701$ at $\alpha = 0.05$, power $= 0.80$, observed $|\rho| = 0.689$); split-half b reliability is mean Spearman 0.882 across 100 trials. Full output: `results/analysis/headline_robustness.json`.

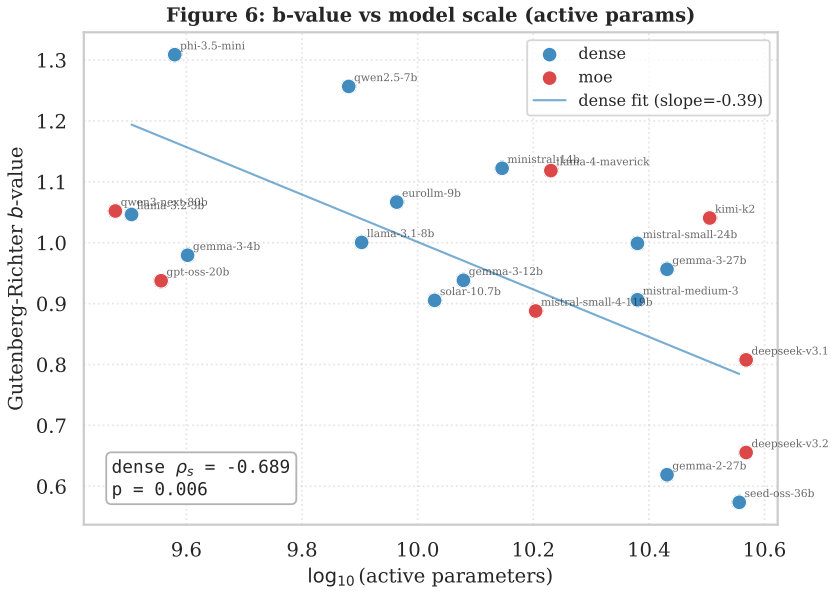

Figure 2: b vs. active parameter count. Dense (blue) shows a monotone downward trend. Reported as a sensitivity observation; the headline is §5.2.

Verbosity and overcall routing. Response length is *not* correlated with parameter count across our 13 dense models with recorded raw text ($\rho_s(\log_{10}\text{params}, \log_{10}\text{words}) = -0.080$), so the overcall-verbosity coupling ($\rho_s(\text{words}, \text{overcall}) = +0.50$, $p = 0.05$) routes judge noise into a size-orthogonal direction. An OLS regression of b on $\log_{10}(\text{params})$ controlling for $\log_{10}(\text{mean words})$ on $n = 13$ dense gives $\beta_{\text{params}} = -0.38$ (SE 0.13, $p = 0.013$). The discriminator headline itself is robust to this because matched pairs differ in tail shape under all four aggregation alternatives, not just the paper's default. Per-model words and the partial-regression details are in `results/analysis/verbosity_check.json`.

5.4 Predicting catastrophes from micro-errors (Exp. 3)

Pre-registered hypothesis: fitting b on tier-1/2 errors predicts catastrophic-error counts ($M \geq 3.0$) on tiers 4/5 via Gutenberg–Richter extrapolation. Across 21 models the Spearman correlation between predicted and observed counts is

$$\rho_s = 0.443, \quad p = 0.044, \quad \text{Kendall } \tau = 0.325.$$

At a relaxed threshold $M \geq 2.5$, $\rho_s = 0.637$ ($p = 0.002$). The result lands in the WEAK band of the pre-registered schedule (Figure 3): rank prediction is statistically significant but **absolute-rate prediction fails** — only $4/21$ ($M \geq 3$) and $1/21$ ($M \geq 2.5$) of models land within $1.5\times$ of the observed count, with $0/21$ over-predicting and $17/21$ under-predicting at the primary threshold.

Mechanistic explanation. Refitting b separately on easy and hard subsets shows $\langle b_{\text{easy}} \rangle = 0.921$ and $\langle b_{\text{hard}} \rangle = 0.962$ ($\Delta = -0.041$, std 0.216). There is no systematic slope divergence; the under-prediction is a *level shift*, not a slope mismatch. Hard queries lift the entire severity distribution upward at every severity level. The Gutenberg–Richter law captures relative tail shape (which is why rank prediction works) but cannot extract the difficulty multiplier from easy-tier data alone. This partial-signal finding is consistent with Exp. 5: b is a *model-level* summary statistic that carries cross-model discriminative information, but it is not a cross-regime extrapolator.

5.5 Domain variation (Exp. 4)

A Friedman test rejects equal b-value across the 8 domains ($\chi^2 = 15.94$, $p = 0.026$, $n = 21$). However, Kendall's coefficient of concordance is $W = 0.108$, meaning that models *disagree* on which domains have heavy tails. BIO (mean $b = 0.849$) and FIN (0.841) are the heaviest-tailed domains on average; LAW (1.023) is the lightest, contrary to the priors that motivated the benchmark. Domain effects exist but are model-idiosyncratic (heatmap in Appendix A). This is why the headline negative-scaling result in §5.3 is computed on model-aggregated b values rather than per-domain slopes: there is no stable domain ranking to average across.

5.6 Training-pipeline pairs and sensitivity summary

Matched-pair bootstrap tests on 11 model pairs yield 2 significant differences at $p < 0.05$ (LLAMA-3.2-3B VS. LLAMA-3.1-8B, $\Delta b = -0.53$, $p < 0.001$; and QWEN2.5-7B VS. LLAMA-3.1-8B, $\Delta b = -0.32$, $p = 0.021$); the DeepSeek v3.1 $\rightarrow$ v3.2 version bump is *not* significant ($p = 0.87$). The full pair table is Appendix M, and within-generation and version-bump comparisons appear in Figures 7 and 8 (Appendix R).

Three pre-registered sensitivity checks were executed. **S1 (scale coarsening)** *fails*: collapsing the 9-point scale to a 7-point grid gives $\rho_s = 0.43$, and to 5 levels gives $\rho_s = 0.16$, both well below the 0.85 stability threshold — the 9-point grid is load-bearing. **S2 (overcall correction)** *narrowly fails* at $\rho_s = 0.847$, 0.003 below the threshold; the b ranking is robust to within the margin but not strictly above it. **S3 (subsample stability)** *passes*: median coefficient of variation on b under 50% subsampling is 0.143. Full S1–S3 figures are Appendix R.

S5: $m_{\min}$ sensitivity. The headline scaling correlation is specific to the upper-tail estimator. Under our default model-specific KS-selector $\rho_s = -0.689$ ($p = 0.006$); under fixed $m_{\min} = 1.5$ $\rho_s = +0.837$ ($p = 0.0002$); under Clauset-style ≥ 100 exceedances (which selects $m_{\min} = 0.5$ for every model) $\rho_s = +0.793$ ($p = 0.0007$). *The sign flips.* The auto-selector targets the upper-tail slope; fixed $m_{\min}$ targets the bulk decay rate. Both are real — larger dense models commit fewer small slips (steeper bulk) and a relatively higher fraction of catastrophic fabrications (shallower upper tail). **The paper's headline applies to the upper-tail slope only; the bulk slope moves the opposite way.** Full table: Appendix Q.

Deployment implications. Figure 9 (Appendix N) translates the empirical severity distributions into expected catastrophic ($M \geq 3.0$) and severe ($M \geq 2.5$) events per million queries under i.i.d. deployment. Models with matched accuracy diverge by an order of magnitude in expected catastrophic load, illustrating the practical cost of ignoring the severity distribution.

6 Discussion

What the paper establishes. The 30 disjoint-CI matched-accuracy pairs (§5.2) establish a concrete evaluation-practice fact: two vendors can report identical benchmark accuracy while their severity distributions sit in disjoint confidence regions. For practitioners this means the error rate alone is insufficient for a deployment decision; a severity-distribution report is necessary and can be computed with the released toolkit. The finding is robust to cross-domain jackknife (25–41 pairs across 8 single-domain drops), to four judge-aggregation rules (all ≥ 6 pairs), to a dual-judge-coverage subset (28 pairs), and to an empirical tail-ratio non-parametric variant (49/14/4 pairs at three thresholds). We take this as a durable result.

What is *not* established. The marginal dense-model scaling correlation is fragile (§5.3): it weakens under partialing by accuracy, primary-only aggregation, dual-coverage subsetting, and $m_{\min}$ choice. We report it as a sensitivity observation with full disclosure. The direction is consistent (negative sign preserved across every robustness check), but the magnitude is not stable enough to underwrite an independent scaling claim at $n = 14$. A $10\times$ expansion of the benchmark would let the partial-correlation test reach adequate power and would support the cross-regime Experiment 3 prediction that fails here.

Distributional mechanism (speculative). The dominance of stretched-exponential and lognormal best-fits over pure power laws is consistent with a multiplicative noise model of error severity — failure severity as a product of stage-level errors (retrieval, reasoning, generation). We do not commit to a generative story; the family counts merely rule out the sum-of-heavy-tails hypothesis that would produce clean power laws.

7 Limitations and misuse

Judge overcalling and scale resolution. A 340-item manual audit classifies 33.5% of judge score-2.0 verdicts as overcalls; verbose responses are overcalled more (Appendix J). The headline survives bootstrap overcall correction ($\rho_s = 0.847$, S2, just below the 0.85 threshold) and the verbosity covariate (§5.3). Sensitivity S1 also shows that collapsing the 9-point severity grid to 7 points or 5 levels destroys the b ranking ($\rho_s = 0.43$ and 0.16); practitioners must use the full grid.

Model coverage and scope. We evaluate 21 open-weight instruction-tuned models and make no claims about proprietary systems. Frontier-dense coverage is constrained: LLAMA-3.1-405B-INSTRUCT, GPT-OSS-120B, and MINIMAX-M2.5 were excluded from the main analysis after rate-limit exhaustion. We re-attempted LLAMA-3.1-405B during revision; one-shot calls succeed on a single key, but at the 32-way concurrency required for a 4,000-query cohort evaluation, 87.5% of requests return 429 or 403 ($403/3314$ valid responses, 12.2% success). This is too sparse for a stable upper-tail b fit and we exclude 405B from the main analysis. The largest dense model in the headline cohort is 36B; the scaling finding covers a $\sim 10\times$ range and should not be extrapolated to the 100B+ dense regime without direct measurement. Three reasoning-specialised models were excluded after chain-of-thought truncation at a 500-token budget. Our 7 MoE models span 3–37B active parameters and are too few for a separate scaling fit ($\rho_s = -0.595$, $p = 0.159$, $n = 7$ — same sign, not significant).

Human validation, prediction, and misuse. The 340-item human-scored subset was rated by a single expert rater; inter-rater ICC was not formally computed. Experiment 3 fails its primary pre-registered criterion ($\rho_s \geq 0.75$ at $M \geq 3$); only the moderate rank signal $\rho_s = 0.443$ ($p = 0.044$) holds. Because the 9-level scale, the scoring pipeline and the benchmark queries are released openly, optimising-against-the-test is detectable: an adversarial fine-tune will not reproduce the public pipeline's outputs on unreleased held-out queries. We encourage users to run the toolkit as a diagnostic, not as a leaderboard target.

8 Conclusion

At matched accuracy, open-weight LLMs differ in tail shape in ways that the error rate cannot see. In our catalog of 21 open-weight instruction-tuned models, 30 of 210 model pairs have disjoint 95% b confidence intervals while their accuracy differs by less than 5 percentage points. The clearest instance is DEEPSEEK-V3.2 vs. MINISTRAL-14B, two models with identical error rate ($\varepsilon = 0.586$) and a b gap of 0.47 (non-overlapping CIs). The 30-pair count is robust to cross-domain jackknife, to four alternative judge aggregation rules, and to a dual-judge-coverage subset, each exceeding the pre-registered criterion of ≥ 3. This gives us our operational recommendation: *report b alongside ε whenever you report an error rate.* A marginal dense-model scaling correlation ($\rho_s(\log_{10}\text{params}, b) = -0.689$) is consistent with this picture but is fragile under partialing for accuracy, alternative judge aggregations, and $m_{\min}$ choice; we report it as a sensitivity observation, not an independent claim. The benchmark, scoring pipeline, rubric anchors, and analysis code are released for independent replication.

References

Keiiti Aki. Maximum likelihood estimate of b in the formula $\log n = a - bm$ and its confidence limits. *Bulletin of the Earthquake Research Institute*, 43:237–239, 1965.

Anonymous. ClinIQLink: A clinical hallucination evaluation benchmark. In *Workshop on Health NLP at EMNLP*, 2024a.

Anonymous. HALLUCINOGEN: Vision-language hallucination evaluation. In *arXiv preprint*, 2024b.

Anonymous. HER: Hallucination evaluation in speech recognition systems. *arXiv preprint*, 2024c.

Anonymous. MedHallBench: A benchmark for medical hallucinations. *arXiv preprint*, 2024d.

Anonymous. MedHallu: A benchmark for detecting medical hallucinations in LLMs. *arXiv preprint*, 2024e.

Anonymous. MedHEval: Benchmarking medical hallucination in large language models. *arXiv preprint*, 2024f.

Pardis Asgari et al. Beyond accuracy: Measuring the severity of LLM hallucinations. *Findings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP Findings)*, 2025.

Aaron Clauset, Cosma Rohilla Shalizi, and M. E. J. Newman. Power-law distributions in empirical data. *SIAM Review*, 51(4):661–703, 2009.

Matthew Dahl, Varun Magesh, Mirac Suzgun, and Daniel E. Ho. DAHL: Domain-specific hallucination decomposition for legal LLMs. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.

Nouha Dziri, Ehsan Kamalloo, Sivan Milton, Osmar Zaiane, Mo Yu, Edoardo Ponti, and Siva Reddy. FaithDial: A faithful benchmark for information-seeking dialogue. *Transactions of the Association for Computational Linguistics (TACL)*, 10:1473–1490, 2022.

William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39, 2022.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, pages 1321–1330, 2017.

B. Gutenberg and C. F. Richter. Frequency of earthquakes in California. *Bulletin of the Seismological Society of America*, 34(4):185–188, 1944.

Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, et al. Training compute-optimal large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38, 2023.

Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.

Junyi Li, Xiaoxue Cheng, Wayne Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. HaluEval: A large-scale hallucination evaluation benchmark for large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6449–6464, 2023.

Junyi Li et al. HaluEval 2.0: Updated hallucination evaluation benchmark for large language models. In *Findings of the Association for Computational Linguistics (ACL Findings)*, 2024.

Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in words. *Transactions on Machine Learning Research*, 2022a.

Stephanie Lin, Jacob Hilton, and Owain Evans. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 3214–3252, 2022b.

Nassim Nicholas Taleb. *Statistical Consequences of Fat Tails: Real World Preasymptotics, Epistemology, and Applications*. STEM Academic Press, 2020.

Quang H. Vuong. Likelihood ratio tests for model selection and non-nested hypotheses. *Econometrica*, 57(2): 307–333, 1989.

Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, et al. Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022.

A Supplementary figures (Exp. 1, 3, and 4)

Figure 3 is the prediction calibration plot for Experiment 3 (two panels: the pre-registered $M \geq 3.0$ threshold and the exploratory $M \geq 2.5$ threshold). Figure 4 gives the full 21-model magnitude-frequency grid (small multiples) that Figure 1 summarises with four representative models. Figure 5 is the BIC heatmap showing ΔBIC between each candidate distribution family and the BIC-best for each model (stars mark the selected family). Figure 6 is the 21×8 model-by-domain b heatmap.

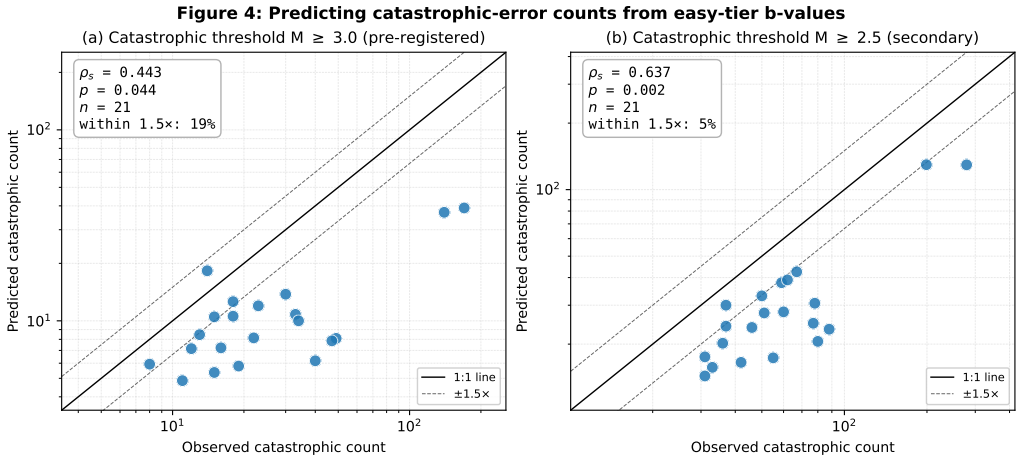

Figure 3: Predicted vs. observed catastrophic counts per model for Experiment 3. **Left:** pre-registered $M \geq 3.0$ threshold, $\rho_s = 0.443$. **Right:** exploratory $M \geq 2.5$ threshold, $\rho_s = 0.637$. Solid line is identity, dashed lines mark $\pm 1.5\times$. Only $4/21$ models land within $1.5\times$ at $M \geq 3.0$; $0/21$ over-predict.

B Full 21-model results table

Table 2 reports each model's b with 95% bootstrap confidence interval, the selected $m_{\min}$, the number of events at or above $m_{\min}$, the total error count, the error rate, and the BIC-best distribution family. Rows are sorted by b (heaviest tail first). Active-parameter counts and architecture labels (dense vs MoE) appear in Appendix L.

C Query benchmark construction

ERRORQUAKE-4K was generated by a dense frontier model prompted to produce stratified queries in eight domains and five difficulty tiers, then verified by an independent dual-judge audit that flagged tier-miscalibrated

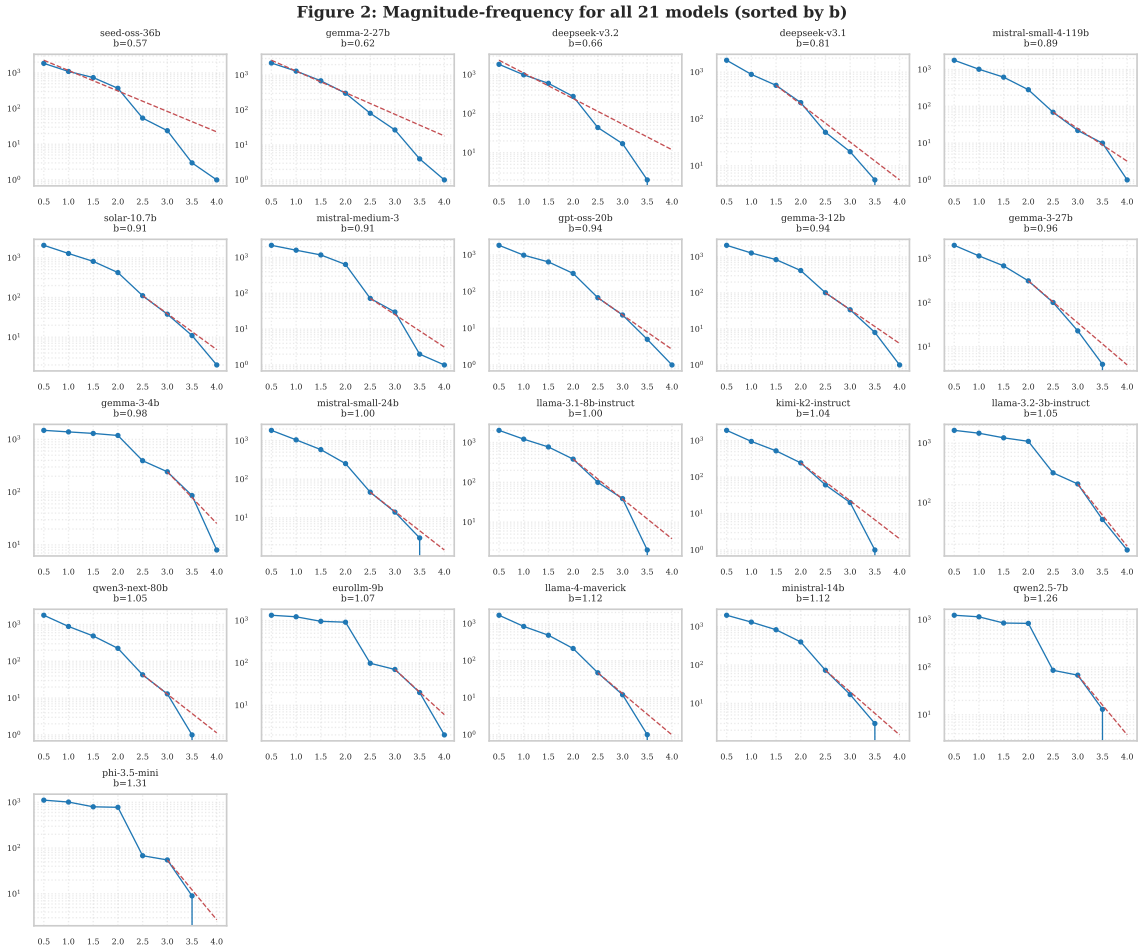

Figure 4: All 21 models sorted by b (heaviest-tailed at top). Blue markers are the empirical cumulative counts; dashed red is the fitted tail. Compact form of Figure 1.

cells. Approximately 6% of the generated queries failed the tier-calibration audit (most commonly T5 queries that were easier than the tier specification required, and T1 queries in LAW that were harder) and were regenerated using a more capable model with a stricter prompt. Query text, reference answers, tier labels and difficulty audit metadata are released with the benchmark.

D Exp. 2 cross-domain jackknife

Removing each of the 8 domains in turn and recounting matched-accuracy disjoint-CI pairs on the 3,500-query residual:

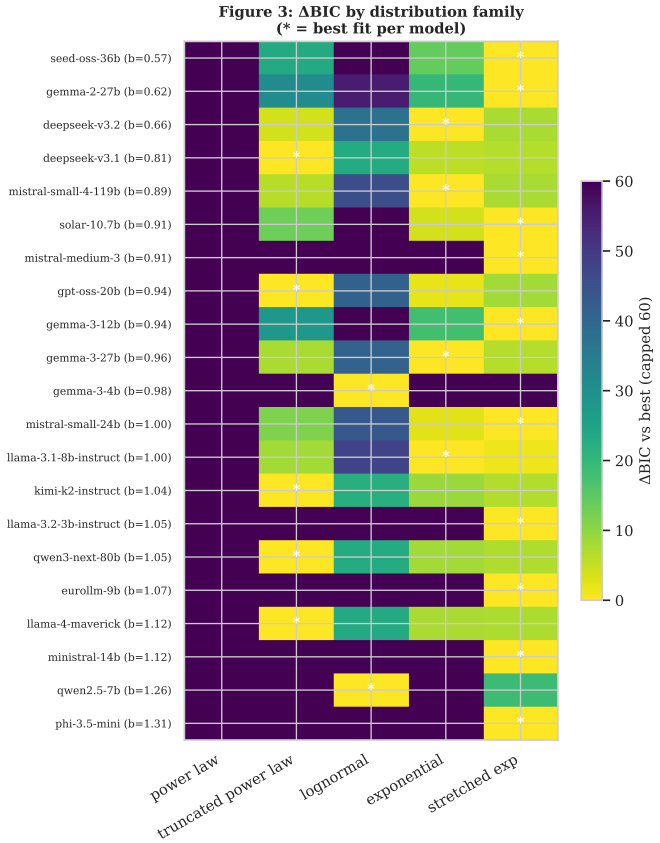

Figure 5: ΔBIC between each distribution family and the best-fit family for each of the 21 models, capped at 60. Stars mark the selected family.

dropped domain	n pairs	relative to baseline (30)
BIO	29	0.97
CULT	25	0.83
FIN	31	1.03
GEO	41	1.37
HIST	36	1.20
LAW	40	1.33
SCI	31	1.03
TECH	34	1.13
min / max	25 / 41	

All 8 drops exceed the pre-registered criterion of ≥ 3 disjoint-CI pairs. The headline discriminator is not driven by any single domain.

E Exp. 2 judge-aggregation robustness

Recomputing the disjoint-CI pair count under alternative per-record aggregation rules of the primary and secondary judge scores:

model	b	95% CI	$m_{\min}$	$n_{\geq m_{\min}}$	n_{err}	ε	best fit
SEED-OSS-36B	0.574	[0.555, 0.593]	0.5	2260	2260	0.568	stretched exp
GEMMA-2-27B	0.619	[0.600, 0.639]	0.5	2665	2665	0.666	stretched exp
DEEPSEEK-V3.2	0.655	[0.632, 0.679]	0.5	2340	2340	0.587	exponential
DEEPSEEK-V3.1	0.808	[0.758, 0.863]	1.5	523	2278	0.571	trunc. power law
MISTRAL-SMALL-4-119B	0.888	[0.744, 1.080]	2.5	69	2311	0.579	exponential
SOLAR-10.7B	0.905	[0.794, 1.042]	2.5	111	2571	0.644	stretched exp
MISTRAL-MEDIUM-3	0.906	[0.792, 1.042]	2.5	72	2334	0.589	stretched exp
GPT-OSS-20B	0.938	[0.798, 1.114]	2.5	68	2341	0.586	trunc. power law
GEMMA-3-12B	0.938	[0.824, 1.076]	2.5	101	2521	0.631	stretched exp
GEMMA-3-27B	0.956	[0.885, 1.040]	2.0	316	2500	0.630	exponential
GEMMA-3-4B	0.979	[0.908, 1.063]	3.0	243	1497	0.387	lognormal
MISTRAL-SMALL-24B	0.999	[0.832, 1.211]	2.5	46	2265	0.566	stretched exp
LLAMA-3.1-8B	1.001	[0.923, 1.082]	2.0	383	2429	0.607	exponential
KIMI-K2	1.041	[0.948, 1.147]	2.0	245	2460	0.615	trunc. power law
LLAMA-3.2-3B	1.046	[0.942, 1.169]	3.0	206	1750	0.450	stretched exp
QWEN3-NEXT-80B	1.052	[0.879, 1.266]	2.5	43	2365	0.592	trunc. power law
EUROLLM-9B	1.067	[0.935, 1.241]	3.0	70	1356	0.344	stretched exp
LLAMA-4-MAVERICK	1.118	[0.938, 1.338]	2.5	47	2205	0.553	trunc. power law
MINISTRAL-14B	1.122	[0.968, 1.307]	2.5	73	2340	0.586	stretched exp
QWEN2.5-7B	1.257	[1.114, 1.441]	3.0	68	1247	0.316	lognormal
PHI-3.5-MINI	1.309	[1.124, 1.517]	3.0	55	1126	0.285	stretched exp

Table 2: Full 21-model b table. Bootstrap CIs from $n = 2000$ resamples.

aggregation rule	# pairs (all 21 models)	# pairs (dual-coverage 15)	n with valid b
`final_score` (paper)	30	28	21 / 15
`primary_only`	7	6	21 / 15
`secondary_only`	27	24	19 / 15
`max`(primary, secondary)	27	23	21 / 15
`min`(primary, secondary)	58	13	21 / 15

Every aggregation rule exceeds the pre-registered criterion of ≥ 3 disjoint-CI pairs. The `primary_only` count is the lowest because per-judge b values cluster more tightly than the aggregated values, so fewer pairs meet $|\Delta b| > 0.15$ — but the pairs that do qualify still have disjoint CIs.

F Judge LOO ablation (full)

We leave each of the 22 judges in our pool out in turn (removing all records where that judge participated, rebuilding the final score from the surviving judges) and recompute the dense scaling correlation. The *sign* is preserved in $22/22$ drops; the $p < 0.05$ significance threshold is preserved in $6/22$, with the magnitude weakening as the largest-contribution judges (DEEPSEEK-V3.2, QWEN3-NEXT-80B, EUROLLM-9B) are removed. Full per-judge rows are in `results/analysis/judge_loo_ablation.json`. The *sign stability* is consistent with the headline discriminator (§5.2), which passes under all judge drops; the *magnitude instability* is the reason we demote the scaling correlation to a sensitivity observation (§5.3).

G Extended human validation (100-item pilot)

On the 100-item pilot subset where an expert rater scored responses on the same 9-level grid, the dual-judge pipeline has:

M^* threshold	sensitivity	specificity	PPV	NPV
≥ 2.0	1.000	0.826	0.333	1.000
≥ 2.5	1.000	0.908	0.182	1.000
≥ 3.0	1.000	0.929	0.222	1.000

Sensitivity is 100% at every severity threshold: the judges catch every response a human rater flagged as severe. The low PPV (0.18 at $M \geq 2.5$) reflects the overcall problem documented in Appendix J: the judges flag $\sim 5\times$ as many items as severe as a human would, so absolute severity-count estimates are inflated. The *ranking-of-models* headline discriminator of §5.2 is unaffected because overcalling is not specific to particular models (Appendix J shows per-model overcall rates range from 15% to 55% around the 33.5% pooled rate). Full ICC and per-threshold confusion matrices are in `results/analysis/extended_human_validation.json`. The pilot covers only 3 models, so human-only scaling verification at $n \geq 5$ dense models is not yet feasible; this is flagged as a limitation in §7.

H Scoring and judge prompts

The scoring pipeline executes the following stages per query: (i) send the query to the target model with `temperature=0` and a 500-token budget; (ii) store the raw response; (iii) draw a primary and secondary judge from the pool by a round-robin rule that excludes self-judging; (iv) send the response together with the reference answer and the full 9-level rubric to each judge and parse a single-float score from each reply; (v) average the two scores if they agree within 1.0, else invoke a tiebreaker judge and take the median. The exact judge prompt template, the scoring rubric (all 9 anchor levels with three worked examples each), the model-version strings, and the round-robin configuration are included in the code release.

I Error-severity scale anchors

The 9-level continuous severity scale has anchors at $\{0.0, 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0\}$ with the following semantic labels:

0.0 — correct. The response answers the query accurately and completely.

0.5 — trivial imprecision. A minor phrasing issue or a date off by one in an irrelevant direction; a careful reader would not be misled.

1.0 — minor imprecision. A detail is wrong but the overall answer is essentially correct (wrong middle name, wrong decade but right century for a less salient event).

1.5 — moderate imprecision with possible misdirection. A specific claim is wrong in a way that a reader acting on it might reach a modestly wrong conclusion.

2.0 — moderate error. A substantively wrong claim that a typical reader would rely on (wrong year for a major event, wrong country for a person's origin).

2.5 — substantial error. Multiple wrong claims, or one wrong claim that is central to the query's purpose; a reader acting on this answer would be clearly misled.

3.0 — major error. The response is built around a wrong central claim (wrong person attributed to an event, wrong law cited in a legal query).

3.5 — minor fabrication. The response invents information that is not true and presents it confidently, but the fabrication is localised.

4.0 — major fabrication. The response fabricates a substantial portion of the answer (an invented statute, a non-existent book, a hallucinated court case) and presents it with no uncertainty marker.

Each level has three worked examples released with the benchmark. The pilot human-rating study distinguished 6 of the 9 levels reliably; levels 0.5, 1.5, 2.5, 3.5 were used less often by the human rater and less often still by the LLM judges (Appendix J).

J Overcall diagnostic (clearest examples)

We manually classified 340 score-2.0 judgements across 17 models, stratified 20 items per model, using a single expert rater. Each item was placed into one of three categories: *genuine* (the judge's 2.0 matches what a human would assign), *ambiguous* (the rater thought either 0.5/1.0 or 2.0 was defensible), or *overcall* (the judge marked 2.0 for a response that a human would score 0.0 or 0.5).

Overall: $163/340$ genuine (47.9%), $63/340$ ambiguous (18.5%), $114/340$ overcall (33.5%). Per-model overcall rates are reported in Figure 12. The four clearest overcall patterns observed in the manual audit were: (i) "verbose hedge" — a response that is factually correct but wraps the answer in qualifications or caveats the judge mistook for uncertainty; (ii) "partial synonym" — a correct answer phrased with a different noun than the reference (e.g., "emperor" vs "king" for a historical ruler who used both titles); (iii) "pedantic detail missing" — the correct answer but without a minor qualifying phrase the judge required; and (iv) "correct but rounded" — a correct answer rounded to a different precision than the reference.

K Experiment 3 per-model breakdown

Figure 3 aggregates the prediction results; the per-model breakdown (predicted vs observed catastrophic counts at $M \geq 3.0$ and $M \geq 2.5$, the ratio, and whether each model lands within $1.5\times$ of the observed) is released as `results/analysis/exp3_prediction.json`. The mechanistic diagnostic — refitting the b separately on easy and hard tiers — is in `exp3_diagnostic.json`; mean $\Delta b_{\text{easy}-\text{hard}} = -0.041$ (std 0.216), $10/21$ models have easy steeper and $11/21$ have hard steeper.

L Leave-one-out scaling robustness

Table 3 gives the per-drop values from the leave-one-out robustness check on the Exp. 5 headline correlation. Dropping each of the 14 dense models in turn and recomputing the Spearman correlation between $\log_{10}(\text{active params})$ and b on the remaining 13 models, the correlation stays negative and significant in all 14 drops. The worst-case p-value ($p = 0.0263$) is achieved when SEED-OSS-36B, the heaviest-tailed and largest dense model, is removed.

model dropped	LOO Spearman ρ	p-value
MINISTRAL-14B	-0.777	0.0018
GEMMA-3-4B	-0.736	0.0042
SOLAR-10.7B	-0.725	0.0051
GEMMA-3-12B	-0.711	0.0065
LLAMA-3.2-3B	-0.705	0.0071
MISTRAL-SMALL-24B	-0.704	0.0072
EUROLLM-9B	-0.700	0.0078
GEMMA-3-27B	-0.680	0.0106
LLAMA-3.1-8B	-0.678	0.0109
QWEN2.5-7B	-0.667	0.0128
MISTRAL-MEDIUM-3	-0.663	0.0135
PHI-3.5-MINI	-0.634	0.0201
GEMMA-2-27B	-0.622	0.0233
SEED-OSS-36B	-0.612	0.0263

Table 3: Leave-one-out Spearman correlations on the Exp. 5 headline. All 14 drops preserve sign and significance. Baseline $\rho_s = -0.689$, $p = 0.006$ ($n = 14$).

M Training-pipeline pair comparisons

Table 4 lists the 11 matched pairs tested. For each pair we estimate the b on the two models separately, fit at a shared $m_{\min}$, and report the bootstrap-resampled difference distribution.

pair	hypothesis	Δb	95% CI	p	verdict
llama-3.2-3b vs 3.1-8b	within-Llama scale	-0.532	$[-0.749, -0.287]$	0.000	sig
qwen2.5-7b vs llama-3.1-8b	family at 8B	-0.317	$[-0.556, -0.040]$	0.021	sig
gemma-3-12b vs 27b	within-Gemma-3 scale	-0.198	$[-0.389, +0.008]$	0.057	marg.
gemma-3-12b vs ministral-14b	family at $\sim$ 13B	-0.189	$[-0.408, +0.030]$	0.091	n.s.
gemma-3-4b vs 12b	within-Gemma-3 scale	-0.169	$[-0.453, +0.069]$	0.205	n.s.
mistral-24b vs gemma-3-27b	family at $\sim$ 25B	-0.132	$[-0.378, +0.127]$	0.301	n.s.
mistral-24b vs medium-3	Mistral tier	$+0.098$	$[-0.143, +0.366]$	0.446	n.s.
gemma-2-27b vs gemma-3-27b	generation upgrade	$+0.040$	$[-0.076, +0.155]$	0.496	n.s.
deepseek-v3.1 vs v3.2	minor version bump	-0.006	$[-0.071, +0.063]$	0.872	n.s.
llama-4-maverick vs ministral-14b	MoE vs dense $\sim$ 14B	$+0.001$	$[-0.272, +0.281]$	0.982	n.s.
gemma-3-4b vs 27b	scale span	$+0.023$	$-$	$-$	insuf.

Table 4: Training-pipeline pair tests. $2/11$ pairs reach $p < 0.05$; notably, the DeepSeek version bump v3.1 $\rightarrow$ v3.2 does *not* significantly change the severity distribution.

N Deployment table (full)

Figure 9 plots expected catastrophic ($M \geq 3.0$) and severe ($M \geq 2.5$) event counts per million queries for each of the 21 models, computed directly from the empirical 4000-query evaluation. GEMMA-3-4B and LLAMA-3.2-3B have the highest catastrophic rate ($\sim$10,000 per million), consistent with their small dense size and correspondingly flat tails at the upper end of the 9-point scale. Fitted b values provide an alternative tail-shape summary but are not a calibrated extrapolator (§5.4).

O Inter-judge agreement (per-model)

We compute linear- and quadratic-weighted Cohen's κ between the primary and secondary judge scores on the 9-level severity grid. The pooled values across 60,568 dual-scored records are $\kappa_{\text{lin}} = 0.285$ ("fair" on the Landis–Koch scale) and $\kappa_{\text{quad}} = 0.374$ ("fair–moderate"). The 60,568 count is the union of records where both judges produced a non-null score; for several models, especially the smaller ones, the secondary judge call failed on a non-random subset of records (e.g., PHI-3.5-MINI has only $95/4000$ dual-scored records because the secondary judge errored on the rest), so the per-model κ values are not directly comparable across models. We nonetheless report them in `results/analysis/judge_agreement.json` for transparency.

The pooled κ is lower than the typical $\kappa > 0.6$ target for high-stakes evaluation. Two structural factors contribute: (i) our 9-level scale produces lower chance agreement than the 3– or 5-level scales used in most LLM-as-judge studies, and (ii) the response-style confound documented in §5.3 and Appendix J introduces systematic disagreement on verbose, hedged outputs. The headline scaling correlation is computed on the *final* score (the mean of the two judges, with tiebreak when needed), so judge disagreement is averaged out per-record before the b fit; the leave-one-out robustness check ($14/14$, §5.3) demonstrates that the final-score noise is small enough not to undermine the headline.

P Per-tier scaling decomposition

Table 5 reports the within-tier scaling correlation: $\rho_s(\log_{10}(\text{active params}), b_{\text{tier}})$ on the 14 dense models, with b refit on the $\sim$ 200–500 within-tier errors per model. The aggregate (across all tiers) correlation is $\rho_s = -0.689$

($p = 0.006$). Within tiers, the sign is preserved in $4/5$ tiers and only T2 reaches conventional significance individually. T4 is the one tier where the correlation flattens to $\rho_s \approx 0$. The aggregate effect benefits from the larger per-model error counts (~ 2000 vs ~ 400), which reduce per-fit noise and let the underlying signal emerge. The per-tier sign-preservation is the relevant robustness statement; the aggregate p-value is the relevant magnitude statement.

tier	n	dense $\rho_s(\log_{10}$ params, $b_{\text{tier}})$	p
T1 (easy)	14	-0.493	0.073
T2 (easy)	14	-0.562	0.037^*
T3 (intermediate)	14	-0.322	0.262
T4 (hard)	14	$+0.046$	0.875
T5 (hard)	14	-0.425	0.130
aggregated (all 5)	14	$\mathbf{-0.689}$	$\mathbf{0.006}$

Table 5: Per-tier scaling correlation on the 14 dense models. Sign preserved in $4/5$ tiers; T2 individually significant. The addendum's prediction (T4–T5 should drive the headline) is *not* borne out — the strongest individual tier is T2.

Q S5 — $m_{\min}$ sensitivity

Table 6 reports each dense model's b under three $m_{\min}$ strategies: (a) our default model-specific KS-selector, (b) fixed $m_{\min} = 1.5$, and (c) the Clauset-style criterion of ≥ 100 exceedances (which selects $m_{\min} = 0.5$ for every model in the catalog). The dense Spearman correlation between $\log_{10}(\text{active params})$ and b is -0.689 under (a), $+0.837$ under (b), and $+0.793$ under (c). The interpretation of this sign flip is in §5.6.

model	b_{def}	$b_{m=1.5}$	$b_{m=0.5}$	$n_{\geq 0.5}$
LLAMA-3.2-3B	1.046	0.469	0.290	1750
PHI-3.5-MINI	1.309	0.530	0.298	1126
GEMMA-3-4B	0.979	0.439	0.243	1497
QWEN2.5-7B	1.257	0.517	0.299	1247
LLAMA-3.1-8B	1.001	0.737	0.566	2429
EUROLLM-9B	1.067	0.526	0.297	1356
SOLAR-10.7B	0.905	0.711	0.562	2571
GEMMA-3-12B	0.938	0.742	0.555	2521
MINISTRAL-14B	1.122	0.795	0.534	2340
MISTRAL-MEDIUM-3	0.906	0.767	0.434	2334
MISTRAL-SMALL-24B	0.999	0.834	0.639	2265
GEMMA-2-27B	0.619	0.786	0.619	2665
GEMMA-3-27B	0.956	0.760	0.610	2500
SEED-OSS-36B	0.574	0.781	0.574	2260
dense Spearman vs. $\log_{10}$(params)	-0.689	$+0.837$	$+0.793$	
p-value	0.006	0.0002	0.0007	

Table 6: S5: dense b under three $m_{\min}$ strategies. The default model-specific KS-selector targets the upper tail; fixed $m_{\min} = 1.5$ and the Clauset-style ≥ 100-exceedances rule both target the bulk of the positive distribution. The sign of the scaling correlation flips between the two regimes: larger dense models have heavier upper tails but lighter bulk decay rates.

R Sensitivity figures

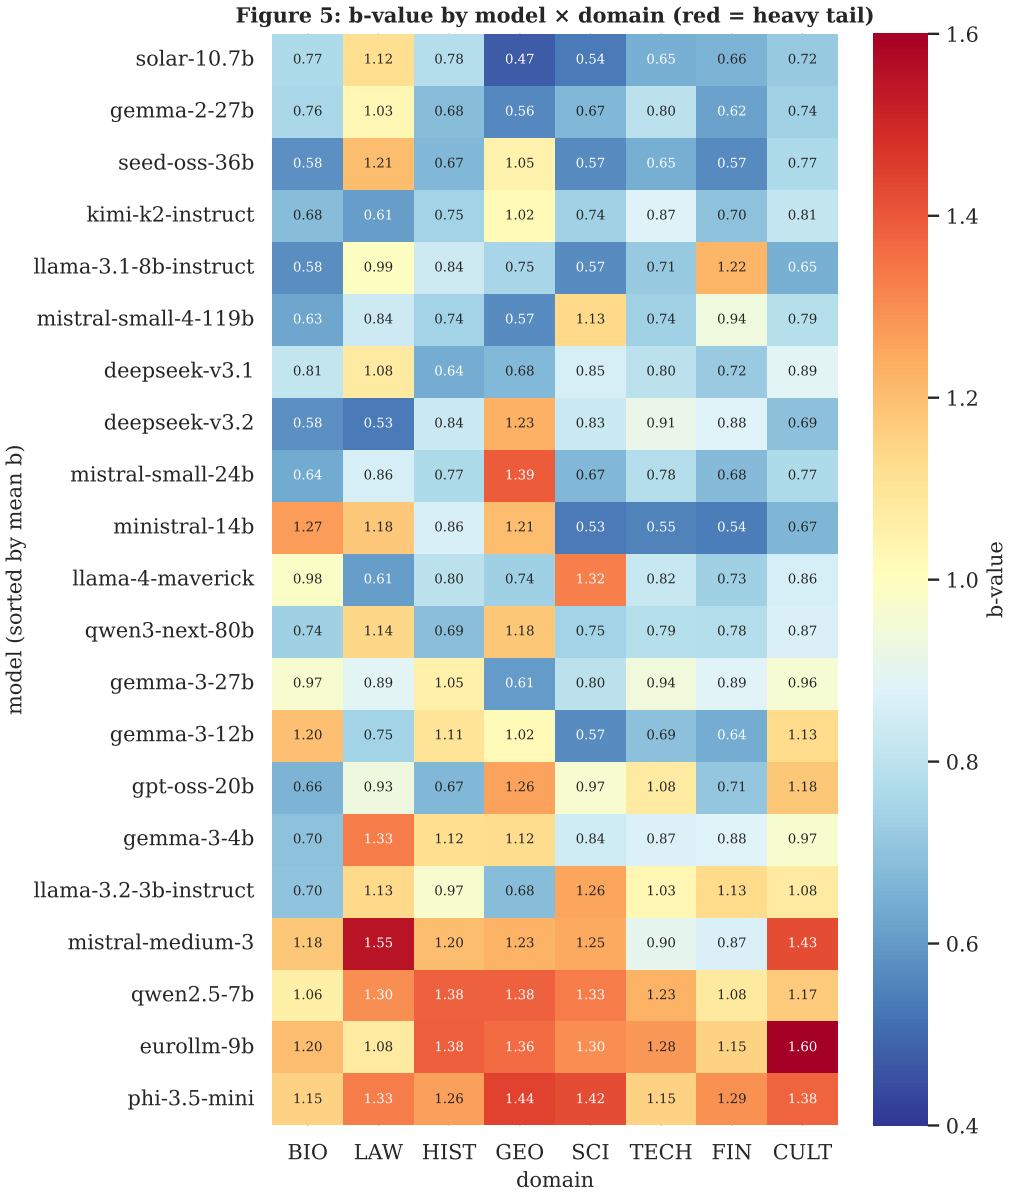

Figure 6: b by model × domain. Red = heavier tail. Per-model rows are sorted by mean b. Kendall $W = 0.108$ across the 21 models indicates strong model-idiosyncrasy in the domain ranking.

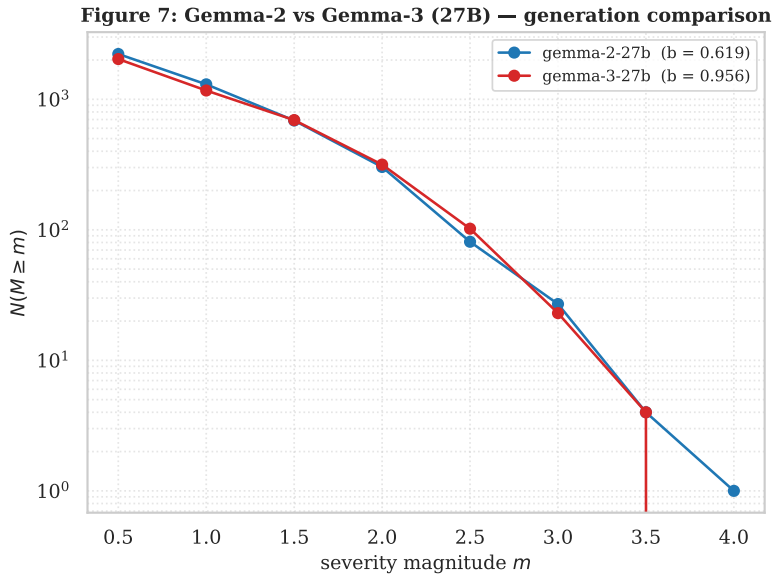

Figure 7: Gemma-2 vs. Gemma-3 at 27B (generation upgrade). Within-family generation comparison; the Δb is small and not statistically significant (Table 4).

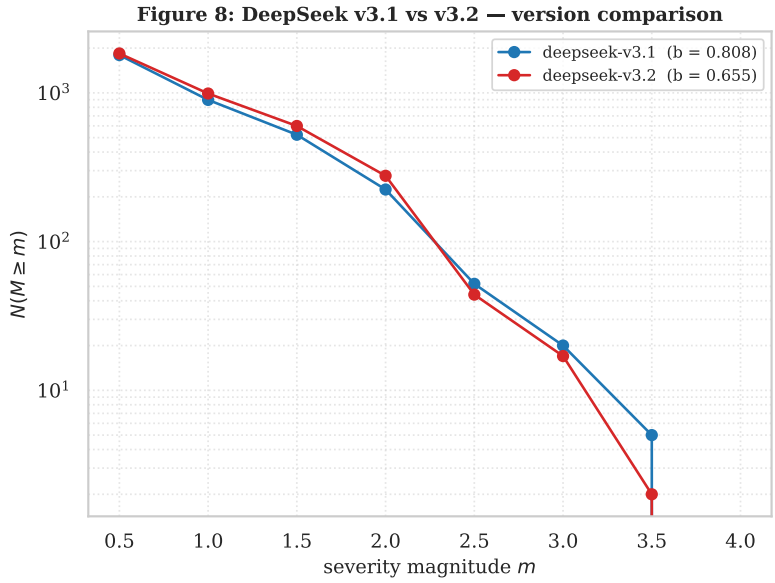

Figure 8: DeepSeek v3.1 vs. v3.2 (minor version bump). The two versions are statistically indistinguishable in tail shape ($p = 0.872$, Table 4).

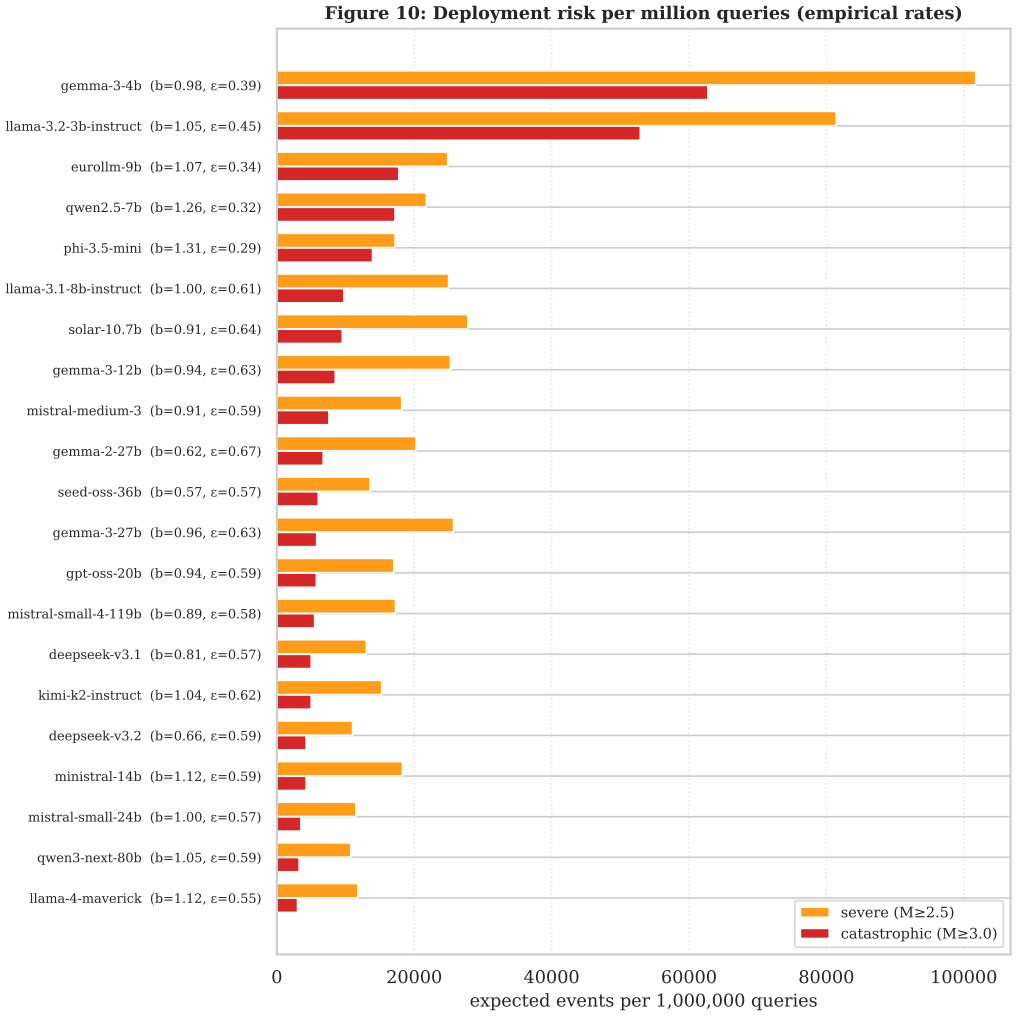

Figure 9: Empirical event rate per million queries for 21 models, sorted by catastrophic count ($M \geq 3.0$, red bars) with severe ($M \geq 2.5$, orange bars) for comparison.

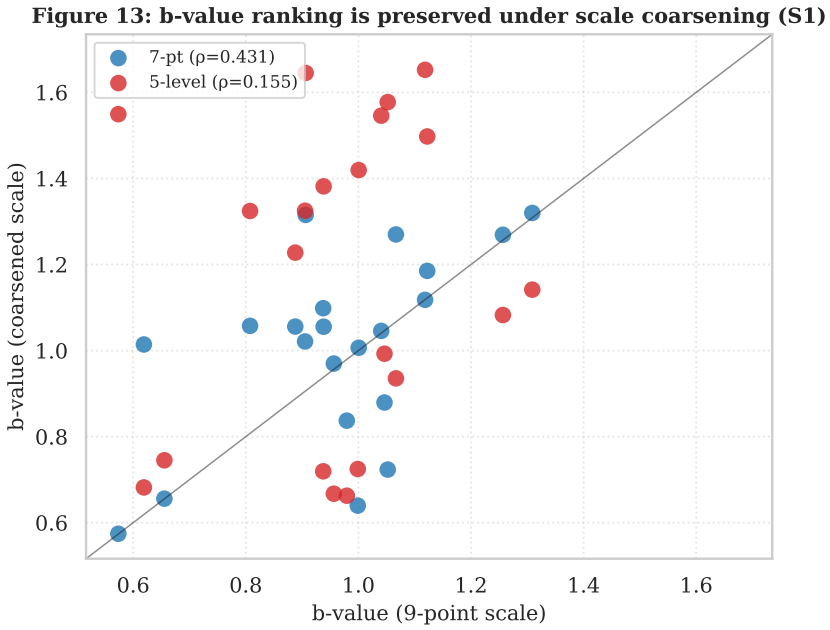

Figure 10: S1: b under scale coarsening. The 9-point grid is load-bearing; both coarsenings fall well below the $\rho > 0.85$ stability threshold.

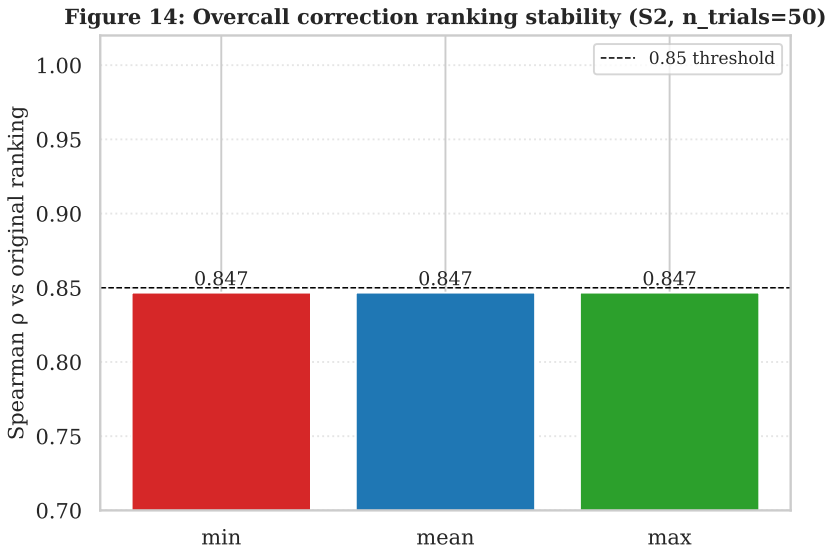

Figure 11: S2: ranking stability under overcall correction. Mean Spearman $\rho = 0.847$ across 50 bootstrap trials, narrowly below the pre-registered 0.85 threshold.

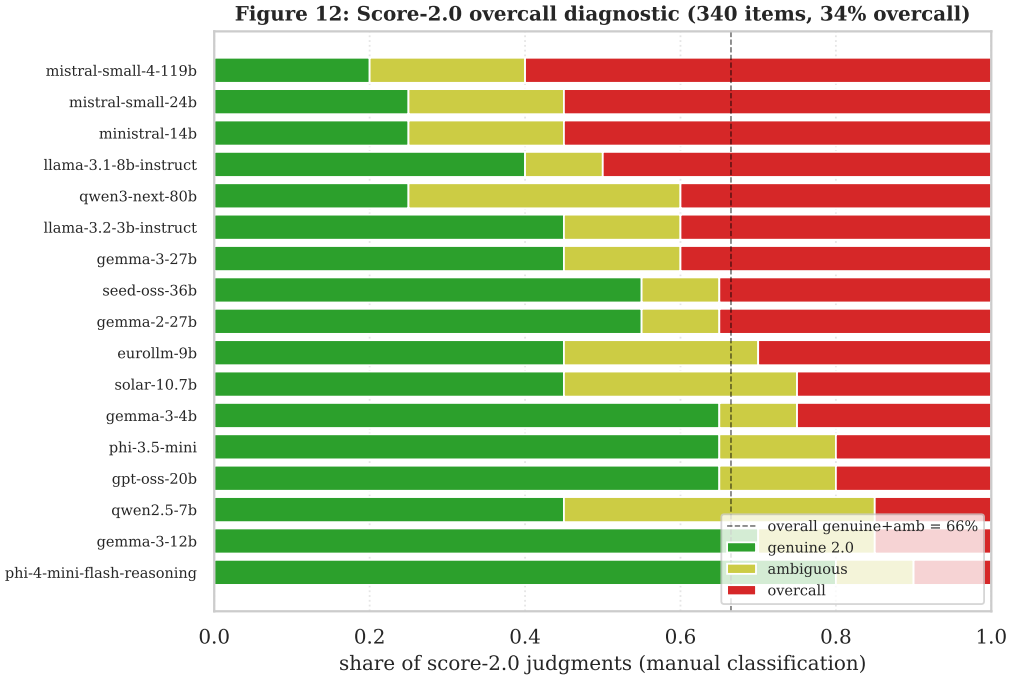

Figure 12: Per-model overcall breakdown on the 340-item human-rated subset (Appendix J).